Spinnaker and the Distributed Monorepo

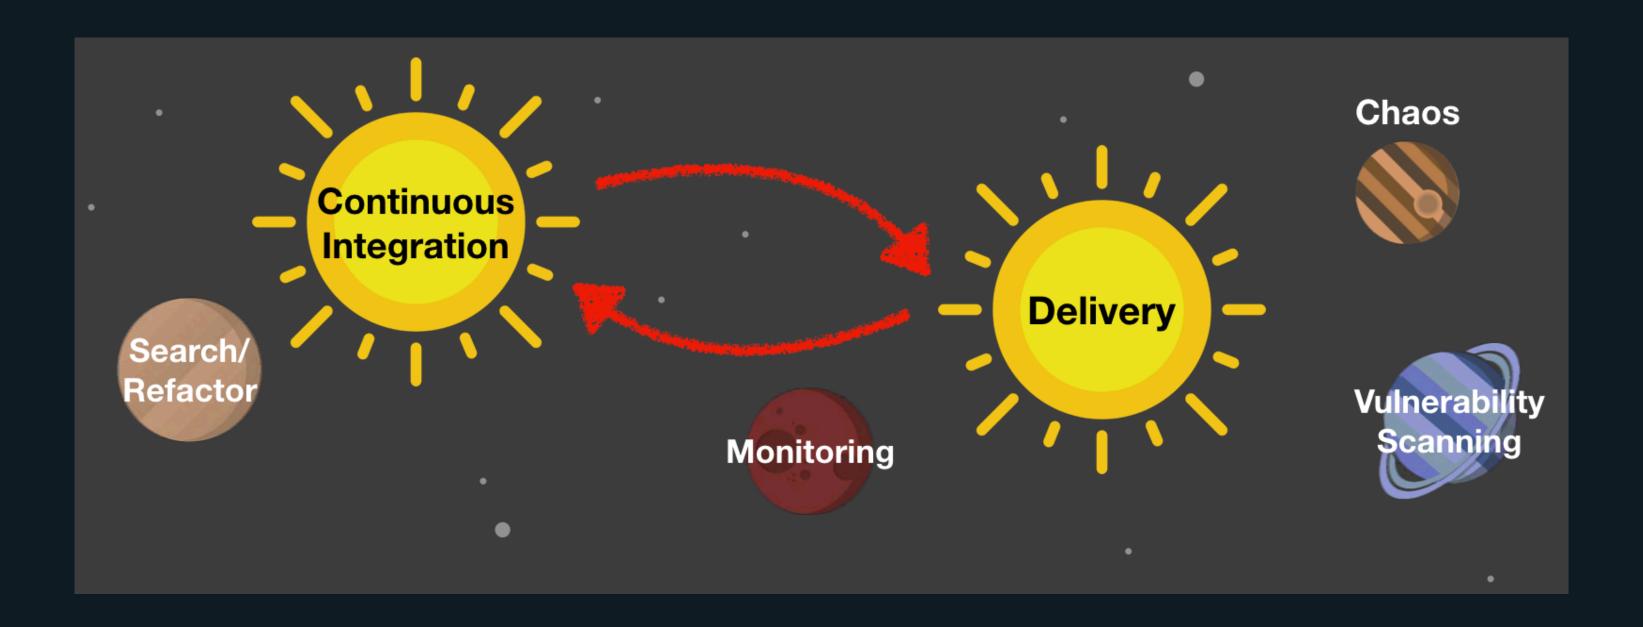
Jon Schneider

@jon_k_schneider

github.com/jkschneider/springone-distributed-monorepo

Part 1: The Spinnaker resource model

CI and CD have distinct orbits

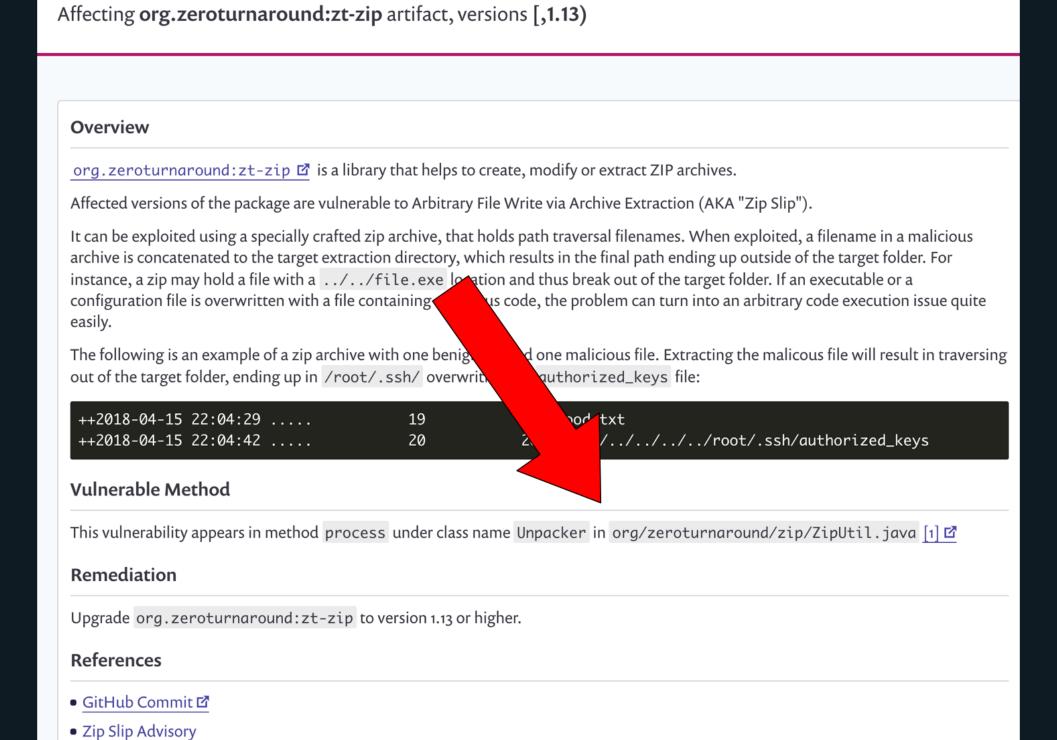


Stitching together source, binaries, and deployed assets.



Ownership service

Part 2: Generate a binary callgraph to search for method level vulnerabilities

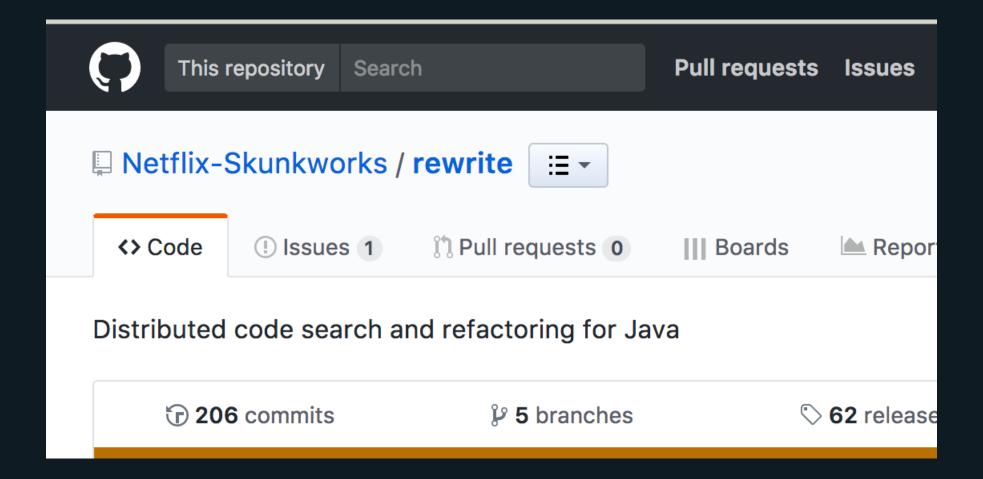


List of fixed projects that contained Zip Slip ☑

6

Part 1: Rewriting code

Rewrite is a programmatic refactoring tool.



Suppose we have a simple class A.

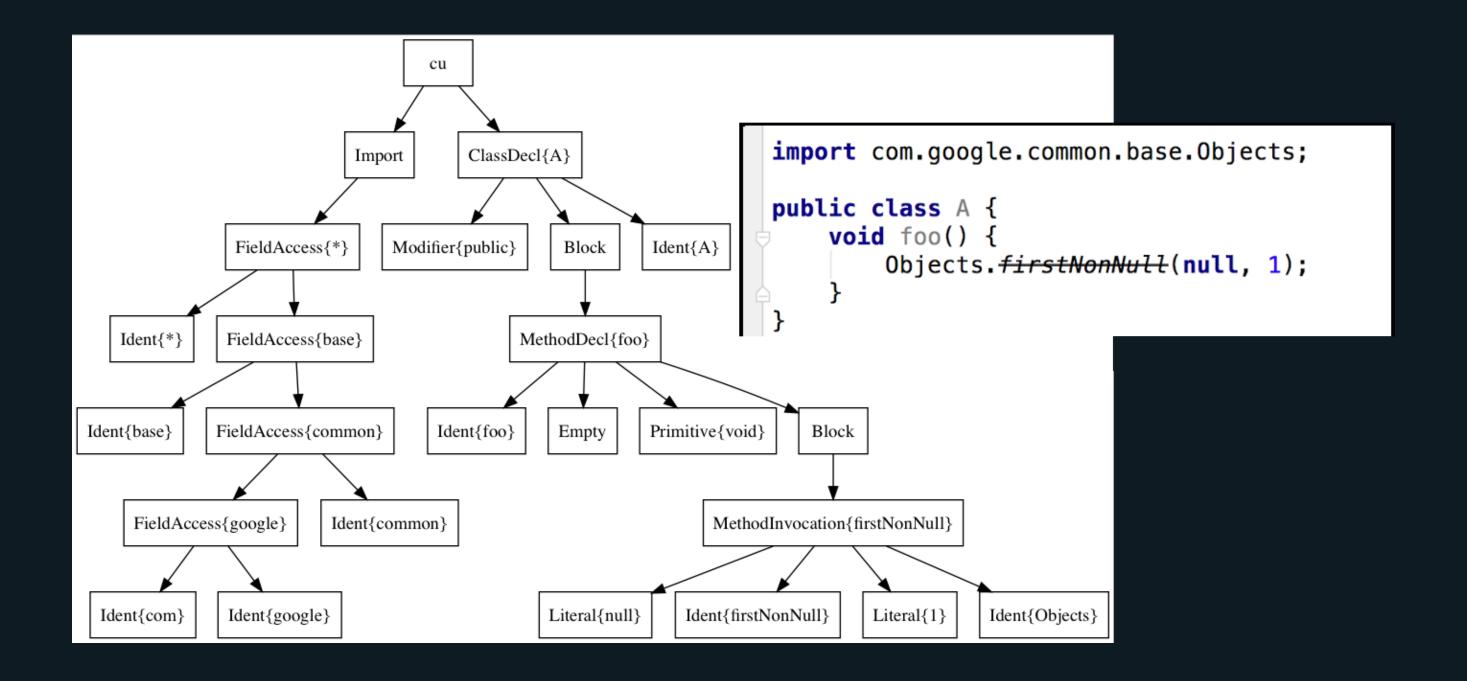
```
import com.google.common.base.Objects;

public class A {
    void foo() {
        Objects.firstNonNull(null, 1);
    }
}
```

9

Raw source code + classpath = Rewrite AST.

The Rewrite AST covers the whole Java language.



Rewrite's AST is special.

- 1. Serializable
- 2. Acyclic
- 3. Type-attributed

Rewrite's AST preserves formatting.

We can find method calls and fields from the AST.

```
Tr.CompilationUnit cu = new OracleJdkParser().parse(aSource);
assertThat(cu.findMethodCalls("java.util.Arrays asList(..)")).hasSize(1);
assertThat(cu.firstClass().findFields("java.util.Arrays")).isEmpty();
```

We can find types from the AST.

```
assertThat(cu.hasType("java.util.Arrays")).isTrue();
assertThat(cu.hasType(Arrays.class)).isTrue();
assertThat(cu.findType(Arrays.class))
   .hasSize(1).hasOnlyElementsOfType(Tr.Ident.class);
```

Suppose we have a class referring to a deprecated Guava method.

We can refactor both deprecated references.

```
Tr.CompilationUnit cu = new OracleJdkParser().parse(bSource);
Refactor refactor = cu.refactor();

refactor.changeMethodTargetToStatic(
   cu.findMethodCalls("com.google..Objects firstNonNull(..)"),
   "com.google.common.base.MoreObjects"
);

refactor.changeMethodName(
   cu.findMethodCalls("com.google..MoreExecutors sameThreadExecutor()"),
   "directExecutor"
);
```

The fixed code emitted from Refactor can be used to overwrite the original source.

// emits a string containing the fixed code, style preserved
refactor.fix().print();

18

Or we can emit a diff that can be used with git apply

// emits a String containing the diff
refactor.diff();

```
diff --git a/B.java b/B.java
index cf08ec7..14f2241 100644
--- a/B.java
+++ b/B.java
@@ -1,15 +1,15 @@
-import com.google.common.base.Objects;
+import com.google.common.base.MoreObjects;
 import com.google.common.util.concurrent.MoreExecutors;
 public class B {
    void foo() {
         Objects.firstNonNull(
        MoreObjects.firstNonNull(
                 null,
                 "hi"
         );
        MoreExecutors.sameThreadExecutor();
        MoreExecutors.directExecutor();
```

```
refactor changeMethodTargetToStatic(
                                                ctor: guava contains all the
                             cu.findMe+hodCalls(signature: "com.google.common.base.Objects firstNonNull(...)"),
refactor.changeMethodTargetToStatic(
                             cu.findMethodCalls( signature: "com.google.common.collect.Iterators emptyIterator(..)"),
                              toClass: "java.util.Collections"
);
refactor.changeMethodName(
                             cu.findMethodCalls(signature: "com.google.common.util.concurrent.MoreExecutors sameThreadEx
                              toName: "directExecutor"
refactor.changeMethodName(
                             cu.findMethodCalls(signature: "com.google.common.util.concurrent.Futures get(java.util.concurrent)
                              toName: "getChecked"
);
                                                                                                                                                                                                                                                                                                                                  20
refactor.changeMethodName(
                            cu.findMethodCalls(signature: "com.google.common.util.concurrent.Futures transform(com.google.common.util.concurrent.Futures transform(com.google.com.google.common.util.concurrent.Futures transform(com.google.common.util.concurrent.Futures transform(com.google.common.util.concurrent.Futures transform(com.google.common.util.concurrent.Futures transform(com.google.common.util.concurrent.Futures transform(com.google.common.util.concurrent.Futures transform(com.google.com.google.com.google.com.google.common.util.com.google.com.google.com.google.com.google.com.google.com.google.com.google.com.google.com.google.com.google.com.google.com.google.com.google.com.google.com.google.com.google.com.
```

Just annotate a static method to define a refactor rule.

Part 2: Using BigQuery to find all Guava code in Github

Identify all Java sources from BigQuery's Github copy.

```
SELECT *
FROM [bigquery-public-data:github_repos.files]
WHERE RIGHT(path, 5) = '.java'
```

23

In options, save the results of this query to:

myproject:springone.java
_files.

You will have to allow large results as well. This is a fairly cheap query (336 GB).

Move Java source file contents to our dataset.

```
SELECT *
FROM [bigquery-public-data:github_repos.contents]
WHERE id IN (
   SELECT id
   FROM [myproject:springone.java_files]
)
Note: This will eat into your $300 credits.
It cost me ~$6 (1.94 TB).
```

Cut down the sources to just those that refer to Guava packages.

Getting cheaper now...

```
SELECT repo_name, path, content
FROM [myproject:springone.java_file_contents] contents
INNER JOIN [myproject:springone.java_files] files
   ON files.id = contents.id
WHERE content CONTAINS 'import com.google.common'
```

2

Notice we are going to join just enough data from springone. java_files and springone: java_file_contents in order to be able to construct our PRs.

Save the result to myproject:springone.java_file_contents_guava.

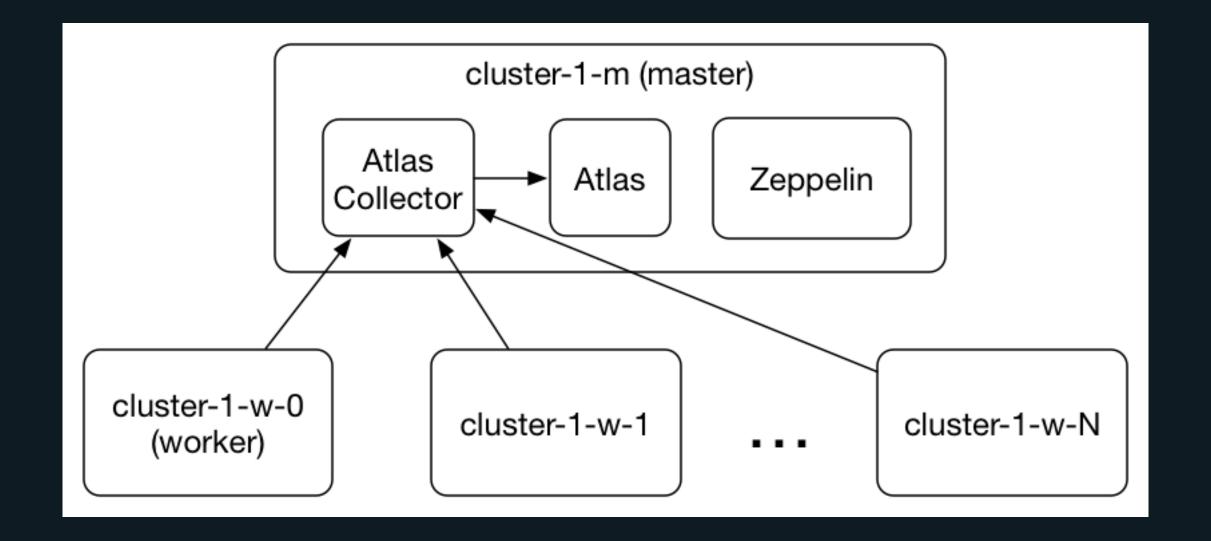
Through Step 3, we have cut down the size of the initial BigQuery public dataset from 1.94 TB to around 25 GB. Much more manageable!

We now have the dataset to run our refactoring rule on.

- 1. 2.6 million Java source files.
- 2. 47,565 Github repositories.

Part 3: Employing our refactoring rule at scale on Google Cloud Dataproc.

Create a Spark/Zeppelin cluster on Google Cloud Dataproc.



Monitoring our Spark workers with Atlas and spring-metrics

```
@SpringBootApplication
@EnableAtlasMetrics // (1)
public class AtlasCollector {
    public static void main(String[] args) {
        SpringApplication.run(AtlasCollector.class, args);
    }
}
```

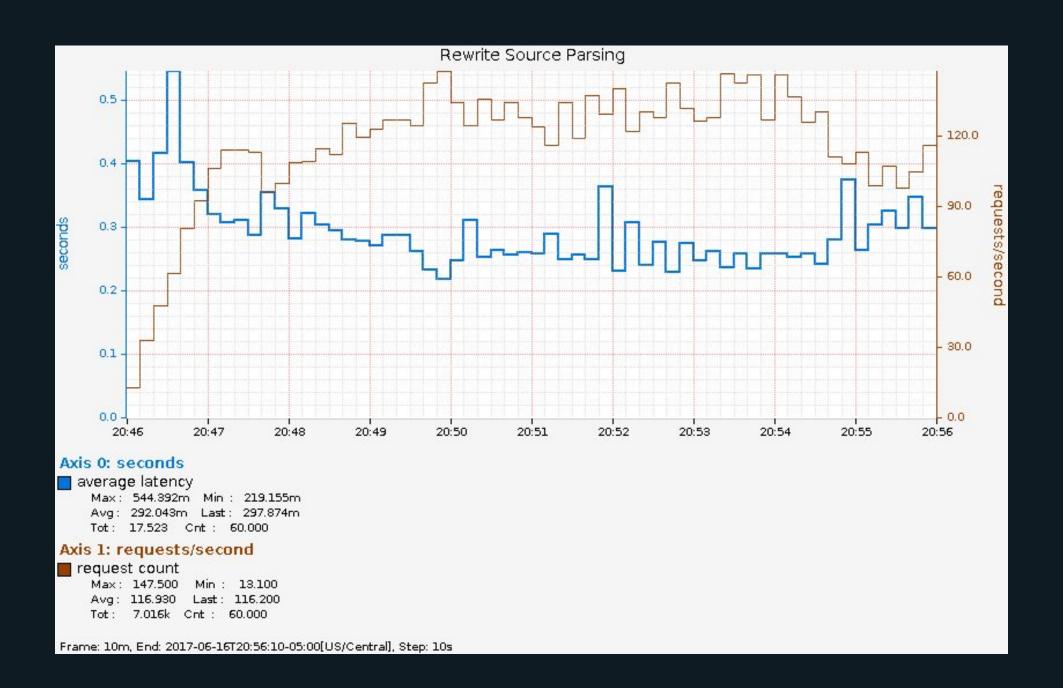
•••

Monitoring our Spark workers with Atlas and spring-metrics

We'll write the job in a Zeppelin notebook.

- 1. Select sources from BigQuery
- 2. Map over all the rows, parsing and running the refactor rule.
- 3. Export our results back to BigQuery.

Measuring our initial pass.



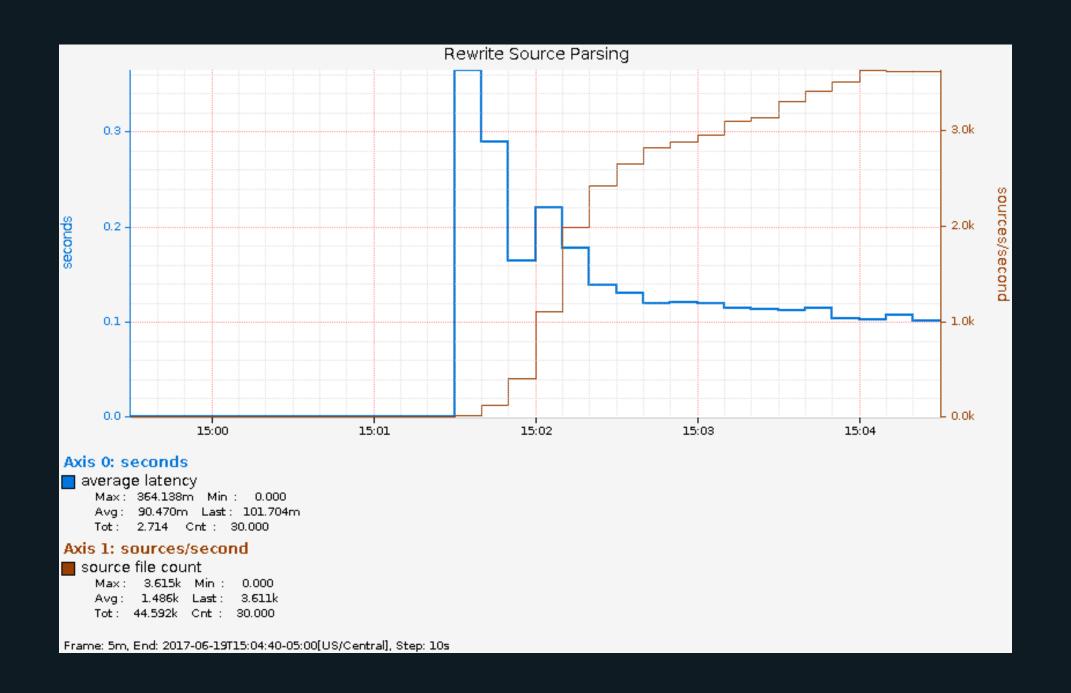
Measuring how big our cluster needs to be.

- 1. Rewrite averages 0.12s per Java source file
- 2. Rate of 6.25 sources per core / second
- 3. With 128 preemptible VMs, we've got: 512 cores * 6.25 sources / core / second

3,200 sources / second = ~13 minutes total

We hope...

After scaling up the cluster with a bunch of cheap VMs.



Some source files are too badly formed to parse.

How many sources did we successfully parse? FINISHED ▷ 💢 🗐 🕸

refactored.count

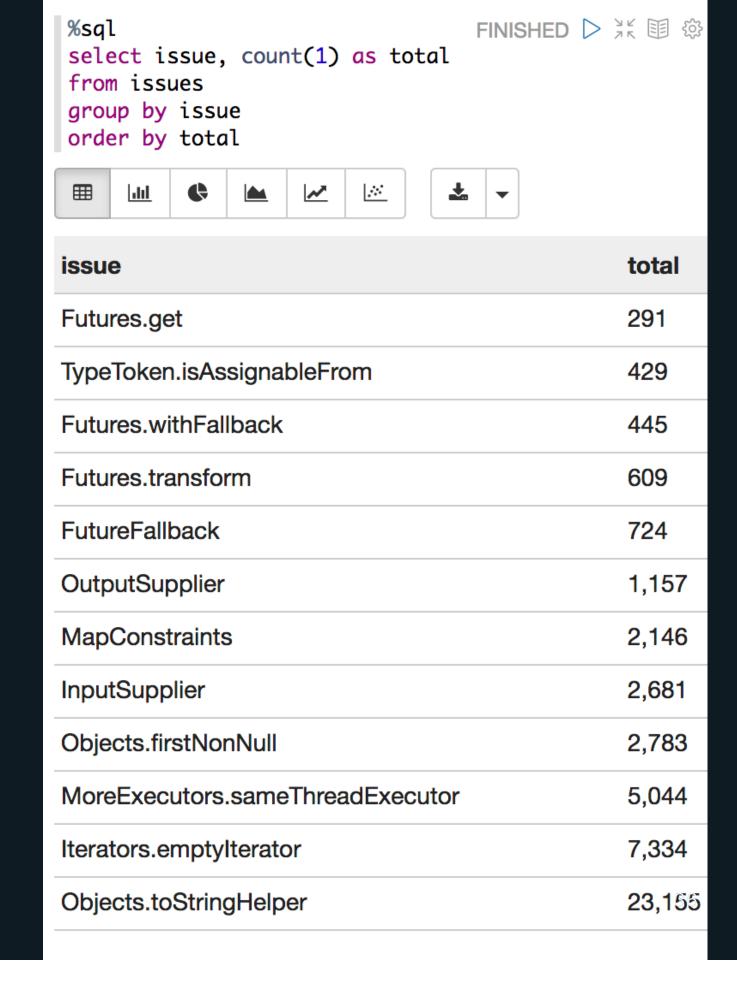
res10: Long = 2590062

Took 10 sec. Last updated by anonymous at June 19 2017, 3:29:40 PM.

2,590,062/2,687,984 Java sources = 96.4%.

We found a healthy number of issues.

- 4,860 of 47,565 projects with problems
- 10.2% of projects with Guava references use deprecated API
- 42,794 source files with problems
- 70,641 lines of code affected



Epilogue: Issuing PRs for all the patches

Generate a single patch file per repo.

SELECT repo, GROUP_CONCAT_UNQUOTED(diff, '\n\n') as patch
FROM [cf-sandbox-jschneider:springone.diffs]
GROUP BY repo

Thanks for attending!